

ML4 – improving weak learners

Machine Learning – Tools and applications for policy – Lecture 5

Iman van Lelyveld – Michiel Nijhuis

DNB Data Science Hub



ML4 – improving weak learners

1. How to grow a decision tree? How to split?
 - Decision trees, purity measures
2. Can Ensemble Classifiers improve weak learners?
 - Bagging, boosting, AdaBoost, XGBoost
3. Can we use forests in other ways?
 - Isolation forests

ML4 – improving weak learners

- Trees

- Forests – Ensemble Classifiers

- Bootstrapping aka Bagging

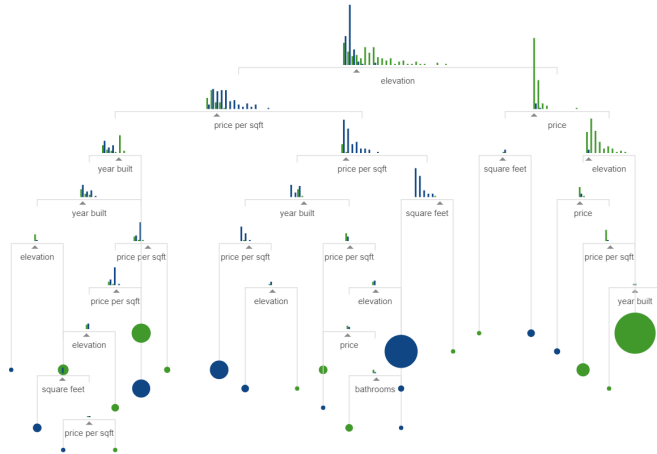
- Boosting

- Isolation forest



- Decision Tree Classifier from Scratch (Josh Gordon) ([link](#))
- ADABOOST clearly explained (Josh Starmer) ([link](#))
- XGBoost from Start to Finish (Josh Starmer) ([link](#))

An extensive webinar where the XGBoost discussion starts at 36.32. The intro is nice to see how to handle data (labeling, missing values, ...)



- A **decision tree** is a learning algorithm that constructs a **set of decisions** based on training data.
- Decision trees are popular because:
 - They are **naturally non-linear**, so you can use them to solve complex problems
 - They are **easy to visualize**
 - How they work is **easily explained**
 - They can be used for **regression** (predict a number) and **classification** (predict a class)
- Drawback: classification at each step only use 'local' information

- At each split we create as much clarity (i.e. information) as possible. The split should make classifying the resulting 'children' easier. Formally (for two-way splits):

$$\text{InformationGain}(D_{\text{parent}}, \text{feature}) = I(D_{\text{par}}) - \frac{N_{\text{left}}}{N_{\text{par}}} I(D_{\text{left}}) - \frac{N_{\text{right}}}{N_{\text{par}}} I(D_{\text{right}})$$

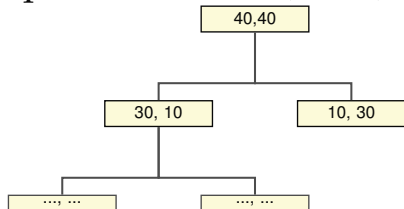
- What could be the Impurity measure?

Gini	$I_{\text{Gini}} = 1 - \sum_{i=1}^c p(i t)^2$
Entropy	$I_{\text{entropy}} = - \sum_{i=1}^c p(i t) \log_2 p(i t)$
Classification error	$I_{\text{classification}} = 1 - \max\{p(i t)\}$


where $p(i|t)$ is the probability that you get it right in 'child' t

- I_{Gini} and I_{entropy} are minimal if homogeneous and maximal if perfectly mixed

Split scenario A (class 0/1)




$$\begin{aligned}
 I_{\text{clas}}(D_{\text{par}}) &= 1 - 0.5 &= 0.5 \\
 I_{\text{clas}}(D_{\text{left}}) &= 1 - \frac{30}{40} &= 0.25 \\
 I_{\text{clas}}(D_{\text{right}}) &= 1 - \frac{30}{40} &= 0.25 \\
 IG_{\text{clas}} &= 0.5 - \frac{40}{80}0.25 - \frac{40}{80}0.25 &= 0.25
 \end{aligned}$$



$$\begin{aligned}
 I_{\text{clas}}(D_{\text{left}}) &= 1 - \frac{40}{60} &= \frac{20}{60} \\
 I_{\text{clas}}(D_{\text{right}}) &= 1 - \frac{20}{20} &= 0 \\
 IG_{\text{clas}} &= 0.5 - \frac{60}{80}\frac{20}{60} - 0 &= 0.25
 \end{aligned}$$


Split scenario B

Split scenario A (class 0/1)



$$\begin{aligned} I_{\text{Gini}}(D_{\text{par}}) &= 1 - (0.5^2 + 0.5^2) &= 0.5 \\ I_{\text{Gini}}(D_{\text{left}}) &= 1 - \left(\frac{10^2}{40} + \frac{30^2}{40}\right) &= \frac{30}{80} \\ I_{\text{Gini}}(D_{\text{right}}) &= 1 - \left(\frac{30^2}{40} + \frac{10^2}{40}\right) &= \frac{30}{80} \\ IG_{\text{Gini}} &= 0.5 - \frac{40}{80} \frac{30}{80} - \frac{40}{80} \frac{30}{80} &= 0.125 \end{aligned}$$


Split scenario B



$$\begin{aligned} I_{\text{Gini}}(D_{\text{left}}) &= 1 - \left(\frac{20^2}{60} + \frac{40^2}{60}\right) &= \frac{4}{9} \\ I_{\text{Gini}}(D_{\text{right}}) &= 1 - (1^2 - 0) &= 0 \\ IG_{\text{Gini}} &= 0.5 - \frac{60}{80} \frac{4}{9} - 0 &= 0.16 \end{aligned}$$



- The **Entropy** criterion also prefers Scenario B over A
- Differences between the I criteria are small
Gini often faster but tends to put most frequent class in 1 branch
- Choice of stopping criteria much more important
 - e.g. until no leaf contains more than x observations
- Overfitting can be reduced by **Cost complexity pruning**: adding a cost to leafy tree. Not implemented in SKLearn library.
- Alternatives:
 - Set maximum depth of tree
 - Set maximum number of leaf nodes
 - Set minimum number of samples per split




Case 1: linear Data Generating Process (DGP)

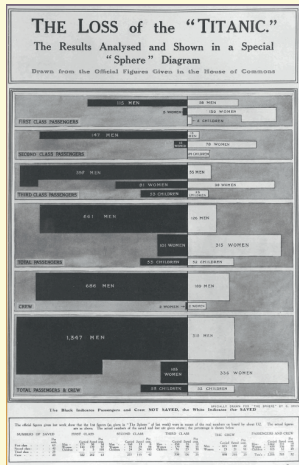
- A linear regression (left) is a better fit than a tree (right).

Case 2: non-linear DGP

- A tree (right) is a better fit than a linear regression (left).



- Following the sinking of the Titanic there was an inquest that documented many features on the passengers
- G.Bron's chart of "The Loss of the 'Titanic'", from The Sphere, 4 May 1912, is the first attempt to show proportions of features (Friendly et al. (2019))
- The data has led to numerous ML flavored analyses (See [here](#), [here](#), [here](#), and [here](#)). These are valuable resources to get to grips with classification trees
- A nice example using temperature data is [here](#)
- See Jake van der Plas's Notebook which can be found [here](#)



ML4 – improving weak learners

Trees

Forests – Ensemble Classifiers

Bootstrapping aka Bagging

Boosting

Isolation forest




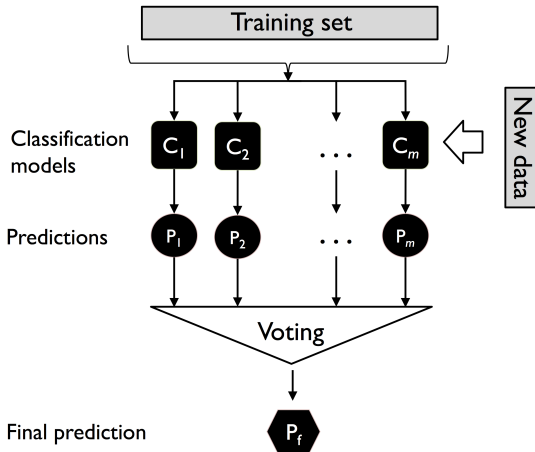
... the interests of truth require a diversity of opinions.

(Mill (1859))

- What if we can learn from the wisdom of crowds? That is, can we **combine multiple classifiers**?
- In a sense we get a “**mixture of experts**”
- Predictions more accurate and robust
- Simplest approach: **majority voting**
- We will discuss the intuition of why this might work

- Majority voting refers to binary setting but can easily generalize to multi-class (e.g. plurality voting)
- Select class label that receives the most votes (**mode**)
- Train m classifiers C_1, \dots, C_m
- Options:
 1. Build ensemble using **different** classification algorithms (e.g. SVM, logistic regression, etc.) but **same** data
 2. Use the **same** algorithm but fit **different subsets** of the training set (e.g. random forest)






We have predictions of individual classifiers C_j and need to select the final class label \hat{y}

$$\hat{y} = \text{mode}\{C_1(\mathbf{x}), C_2(\mathbf{x}), \dots, C_m(\mathbf{x})\}$$

For example, in a binary classification task where $\text{class}_1 = -1$ and $\text{class}_2 = +1$, we can write the majority vote prediction as follows:

$$C(\mathbf{x}) = \text{sign}\left[\sum_j^m C_j(\mathbf{x})\right] = \begin{cases} 1 & \text{if } \sum_j C_j(\mathbf{x}) \geq 0 \\ -1 & \text{otherwise} \end{cases}$$



Rohith Ghandi

- If errors ϵ are more or less independent a **combination** will **reduce the variance**
- Assume that all n base classifiers have the same – but independent – error rate ϵ
- The probability of an error of an ensemble can be expressed as a probability mass function of a **binomial distribution**:

$$P(y \geq k) = \sum_k^n \binom{n}{k} \epsilon^k (1 - \epsilon)^{n-k} = \epsilon_{\text{ensemble}}$$

- Here, $\binom{n}{k}$ is the binomial coefficient and the distribution computes what the probability is of k successes in n trials. In other words, we compute the **probability** that the **prediction of the ensemble is wrong**.


Imagine we have 11 base classifiers ($n = 11$) with individual error rates ϵ of 0.25 and:

$$P(y \geq k) = \sum_{k=6}^{11} \binom{11}{k} 0.25^k (1 - 0.25)^{11-k} = 0.034$$

So the error rate of the ensemble of $n = 11$ classifiers (0.034) is much lower than the error rate of the individual classifiers (0.25).

Same reasoning applied to a wider range of error rates

20



- Simply instantiate several classifiers and make a list
- Pass to `sklearn.ensemble.VotingClassifier(...)`

```
1 clf1 = LogisticRegression(random_state=1)
2 clf2 = RandomForestClassifier(random_state=1)
3 clf3 = GaussianNB()
4 estimators=[('lr', clf1), ('rf', clf2), ('gnb', clf3)]
5 ens_clf = VotingClassifier(estimators)
6 ens_clf = ens_clf.fit(X, y)
```

- If all classifiers can estimate class probabilities, then using **soft voting** might help. In contrast with **hard voting**, confident estimators are given more weight

► [sklearn API link](#)

ML4 – improving weak learners

Trees

Forests – Ensemble Classifiers

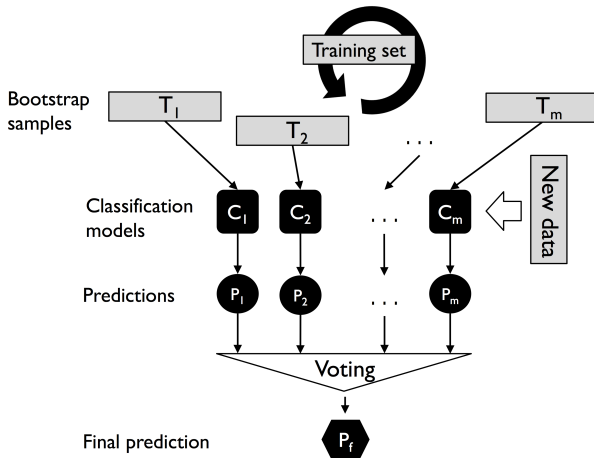
Bootstrapping aka Bagging

Boosting

Isolation forest

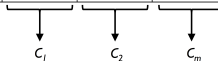


- We used the entire training set for the majority vote classifier
- Here we draw random **bootstrap samples**
 - In statistics, **bootstrapping** is any test or metric that relies on **random sampling with replacement**
 - bootstrapping often used as an alternative to statistical inference based on the assumption of a parametric model when that assumption is in doubt
 - Alternatively we can sample *without* replacement: **pasting**
- The basic idea of bootstrapping is that inference about a population from sample data, can be modeled by resampling with replacement the sample data and performing inference about a sample from resampled data



- Seven training examples
- Sample randomly with replacement
- Use each bootstrap sample to train a classifier C_j
- C_j is typically a decision tree
- **Random Forests**: also use random feature subsets
 - Instantiate a decision tree classifier
 - Make a bagging classifier with decision trees
 - Check that the accuracy is higher for the bagging classifier
- Asses model on **out-of-bag (oob)** observations (oob_score=True)
- Scales nicely because each of the models can be trained separately

Sample indices	Bagging round 1	Bagging round 2	...
1	2	7	...
2	2	3	...
3	1	2	...
4	3	1	...
5	7	1	...
6	2	7	...
7	4	7	...



ML4 – improving weak learners

Trees

Forests – Ensemble Classifiers

Bootstrapping aka Bagging

Boosting

Isolation forest




- Basic idea: start with weak learners that have only a slight performance advantage over random guessing (e.g. a decision tree 'stump') and try to boost their performance by **focusing on training samples that are hard to classify**
- Very simple base classifiers learn from misclassified training examples
- The original boosting algorithm was formulated by Robert Schapire in 1990
- It was later refined into **AdaBoost** – short for Adaptive Boosting
- **AdaBoost** used to be the common implementation of boosting. Lately **XGBoost** is becoming very popular.


1. Draw a **random subset** of training samples d_1 **without replacement** from the training set D to train a weak learner C_1
2. Draw second random training subset d_2 without replacement from the training set and add 50 percent of the samples that were previously misclassified to train a weak learner C_2
3. Find the training samples d_3 in the training set D on which C_1 and C_2 disagree to train a third weak learner C_3
4. Combine the weak learners C_1 , C_2 , and C_3 via majority voting

- In contrast, AdaBoost uses the complete training set to train the weak learners
- Training samples are re-weighted in each iteration to build a strong classifier
- End goal is to build a strong classifier that learns from the mistakes of the previous weak learners in the ensemble


Learner D1



Learner




- In contrast, AdaBoost uses the complete training set to train the weak learners
- Training samples are re-weighted in each iteration to build a strong classifier
- End goal is to build a strong classifier that learns from the mistakes of the previous weak learners in the ensemble




```
1      # AdaBoost Algorithm
2      from sklearn.ensemble import
          AdaBoostClassifier
          AdaBoostClassifier()
3      # n_estimators = 50 (default value)
4      # base_estimator = DecisionTreeClassifier
          (default value)
5      clf.fit(x_train, y_train)
6      clf.predict(x_test)
```

- View boosting problem as an optimization problem, i.e. we take up a loss function and try to optimize it
- New **weak learners** are **added** to concentrate on the areas **where the existing learners** are **doing poorly**

Iteration 1



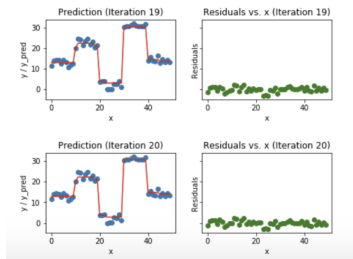
Iteration 2



... last round

- View boosting problem as an optimization problem, i.e. we take up a loss function and try to optimize it
- New weak learners are added to concentrate on the areas where the existing learners are performing poorly

```
1      # Gradient Boosting
2      from sklearn.ensemble import
          GradientBoostingClassifier
3      clf = GradientBoostingClassifier()
4      # n_estimators = 100 (default)
5      # loss function = deviance(default) used
          in Logistic Regression
6      clf.fit(x_train, y_train)
7      clf.predict(x_test)
```



- XGBoost is similar to gradient boosting algorithm, but it has a few smart features
 - Clever penalization of Trees
 - A proportional shrinking of leaf nodes
 - Newton Boosting
 - Extra randomization parameter to reduce correlation between learners

```
1  # XGBoost
2  from xgboost import XGBClassifier
3  clf = XGBClassifier()
4  # n_estimators = 100 (default)
5  # max_depth = 3 (default)
6  clf.fit(x_train, y_train)
7  clf.predict(x_test)
```

- See the [in-depth Webinar](#) by Josh Starmer for a walk through. XGBoost starts at 36.32 but the intro is nice to see how to handle data.

- Ensemble learning is also used a lot in finance/econometric forecasting
 - Bates and Granger (1969), Rapach, Strauss and Zhou (2010), Stock and Watson (2004)
- Puzzle: Many, even sophisticated, models cannot beat a simple mean combination of univariate stock market return forecasts (Jacobsen et al. 2021)
- Especially challenging in streaming applications (Boulegane 2021)

ML4 – improving weak learners

Trees

Forests – Ensemble Classifiers

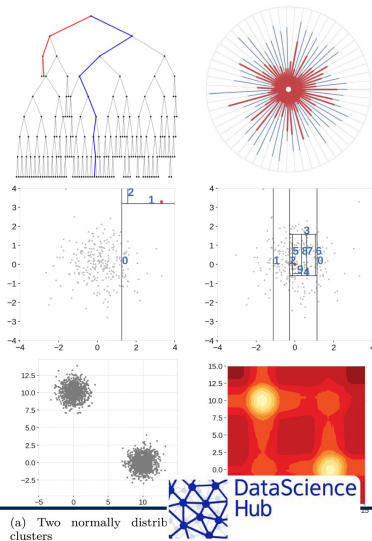
Bootstrapping aka Bagging

Boosting


Isolation forest

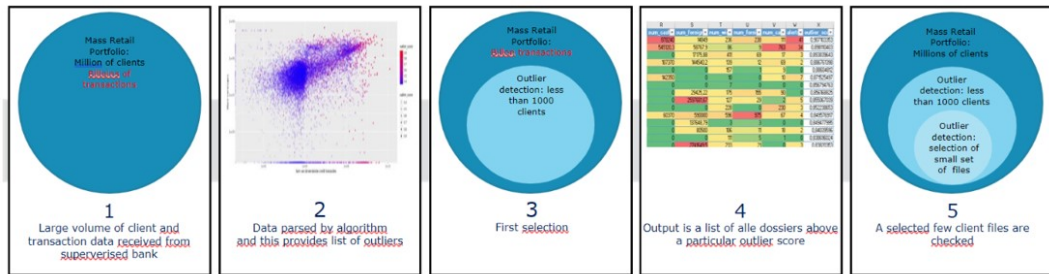


- 'Isolates' observations by **randomly selecting a feature** and then **randomly selecting a split value** ($\min. < \text{split} < \max.$)
- Since recursive partitioning can be represented by a tree structure, the number of splits required to isolate a sample equals the path length from the root node to the terminating node = **Shorter path length**, the more likely it is that it is an **anomaly**
- Can be expressed as an 'outlier score'





- Extended isolation forest selects **random intercept** and **random slope** (Hariri et al. (2021))
- Needs to store less information → faster, especially with many dimensions





- **Stacking** combines the outputs from multiple base models into a single score. The base-level models are trained based on a complete dataset, and then their outputs are used as input features to train an ensemble function
- Usually, the ensemble function is a simple linear combination of the base model scores.
- **Deep stacking** feeds the ensemble outcome(s) into a neural net + the original data



1. Split the data into a training and validation set
2. Divide the training set into K folds, for example 10
3. Train a base model (say SVM) on 9 folds and make predictions on the 10th fold
4. Repeat until you have a prediction for each fold
5. Fit the base model on the whole training set
6. Use the model to make predictions on the test set
7. Repeat step 3 – 6 for other base models (for example decision trees)
8. Use predictions from the test set as features to a new model – the meta-model
9. Make final predictions on the test set using the meta model

With regression problems, the values passed to the meta-model are numeric. With classification problems, they are probabilities or class labels.

- Blending \cong stacking, but uses a holdout set from the training set to make predictions. Predictions are done on the holdout set only. Predictions and holdout set are used to build a final model for predictions on the test set. Blending can be seen as stacking, where the meta-model is trained on predictions made by the base model on the hold-out validation set.
- The term blending was introduced by the Netflix Prize in Kaggle.
 - The blended solution achieves a 10-fold performance improvement on Netflix's movie recommendation algorithm.






- Steps:
 1. Split the data into a test and validation set
 2. Fit base models on the validation set
 3. Make predictions on the validation and test set
 4. Use the validation set and its predictions to build a final model
 5. Make final predictions using this model
- Blending vs stacking:
 - Blending is simpler than stacking and prevents leakage of information in the model. The generalizers and the stackers use different datasets. However, blending uses less data and may lead to overfitting.
 - Cross-validation is more solid on stacking than blending. It is calculated over more folds, compared to using a small hold-out dataset in blending.

In this lecture we covered:

1. How to define a tree, specifically how to split at each node
2. How to combine classifiers: ensemble learning a.k.a. “forests”
3. How ensembles can combine weak learners to come to much stronger learners. In particular AdaBoost and XGBoost
4. How we can use isolation forests for outlier detection



1. Set weight vector \mathbf{w} to uniform weights where $\sum_i w_i = 1$
2. For j in m boosting rounds, do the following:
 1. Train a weighted weak learner $C_j = \text{train}(\mathbf{X}, \mathbf{y}, \mathbf{w})$
 2. Predict class labels $\hat{\mathbf{y}} = \text{predict}(C_j, \mathbf{X})$
 3. Compute the weighted error rate $\epsilon = \mathbf{w} \cdot (\hat{\mathbf{y}} \neq \mathbf{y})$
 4. Compute the coefficient $\alpha_j = 0.5 \log \frac{1-\epsilon}{\epsilon}$
 5. Update the weights $\mathbf{w} := \mathbf{w} \times \exp(-\alpha_j \times \hat{\mathbf{y}} \times \mathbf{y})$
 6. Normalize weights to sum to 1 $\mathbf{w} := \mathbf{w} / \sum_i w_i$
3. Compute the final prediction: $\hat{\mathbf{y}} = (\sum_{j=1}^m (\alpha_j \times \text{predict}(C_j, \mathbf{X})) > 0)$

-
-  Boulegane, D. (2021). Machine learning algorithms for dynamic Internet of Things. PhD thesis.
 -  Friendly, M., Symanzik, J., & Onder, O. (2019). Visualising the Titanic disaster. Significance, 16(1), 14–19.
 -  Hariri, S., Kind, M. C., & Brunner, R. J. (2021). Extended Isolation Forest. IEEE Transactions on Knowledge and Data Engineering, 33(4), 1479–1489.
 -  Jacobsen, B., Jiang, F., & Zhang, H. (2021). Ensemble Machine Learning and Stock Return Predictability. Working Paper.
 -  Mill, J. S. (1859). On Liberty. Cambridge University Press.